

## 1. Mount your Google Drive in Colab

```
from google.colab import drive
drive.mount('/content/drive')
```

↻ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## 2. Access the dataset from your Google Drive

```
import pandas as pd
file_path = '/content/drive/MyDrive/fraud_det/historical_dataset.csv'
df = pd.read_csv(file_path)
```

## 3. Basic exploration to understand the data

```
# Show basic info
print("🔍 Dataset Info:")
print(df.info())

# Check for missing values
print("\n💡 Null Values:")
print(df.isnull().sum())

# Show class distribution
print("\n📊 Class Distribution (isFraud):")
print(df['isFraud'].value_counts(normalize=True))

# Preview the dataset
print("\n👁 First 5 Rows:")
print(df.head())
```

↻ 🔍 Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6362620 entries, 0 to 6362619
Data columns (total 11 columns):
 #   Column                Dtype
---  -
 0   step                  int64
 1   type                  object
 2   amount               float64
 3   nameOrig              object
 4   oldbalanceOrg         float64
 5   newbalanceOrig        float64
 6   nameDest              object
 7   oldbalanceDest        float64
 8   newbalanceDest        float64
 9   isFraud               int64
10   isFlaggedFraud        int64
dtypes: float64(5), int64(3), object(3)
memory usage: 534.0+ MB
None
```

## 💡 Null Values:

```
step          0
type          0
amount        0
nameOrig      0
oldbalanceOrg 0
newbalanceOrig 0
nameDest      0
oldbalanceDest 0
newbalanceDest 0
isFraud       0
isFlaggedFraud 0
dtype: int64
```

## 📊 Class Distribution (isFraud):

```
isFraud
0    0.998709
1    0.001291
Name: proportion, dtype: float64
```

## 👁 First 5 Rows:

```
   step  type  amount  nameOrig  oldbalanceOrg  newbalanceOrig  \
0     1  PAYMENT  9839.64  C1231006815    170136.0    160296.36
1     1  PAYMENT  1864.28  C1666544295     21249.0    19384.72
2     1  TRANSFER   181.00  C1305486145        181.0         0.00
3     1  CASH_OUT   181.00  C840083671        181.0         0.00
```

4	1	PAYMENT	11668.14	C2048537720	41554.0	29885.86
		nameDest	oldbalanceDest	newbalanceDest	isFraud	isFlaggedFraud
0	M1979787155		0.0	0.0	0	0
1	M2044282225		0.0	0.0	0	0
2	C553264065		0.0	0.0	1	0
3	C38997010		21182.0	0.0	1	0
4	M1230701703		0.0	0.0	0	0

#### 4. Handle missing values

```
df.loc[:, 'amount'] = df['amount'].fillna(df['amount'].mean())
print(df.isnull().sum()) # Should show 0 missing values
```

```
step      0
type      0
amount    0
nameOrig  0
oldbalanceOrig  0
newbalanceOrig  0
nameDest  0
oldbalanceDest  0
newbalanceDest  0
isFraud    0
isFlaggedFraud  0
dtype: int64
```

#### 5. Data preprocessing :

- Drop unnecessary columns (nameOrig, nameDest) — these are IDs and not useful for prediction.
- Convert the type column (categorical) into numerical using one-hot encoding.
- Prepare the feature matrix X and the target y.

```
# Drop ID columns
df_clean = df.drop(['nameOrig', 'nameDest'], axis=1)

# One-hot encode the 'type' column
df_clean = pd.get_dummies(df_clean, columns=['type'])

# Separate features and target
X = df_clean.drop(['isFraud'], axis=1)
y = df_clean['isFraud']

print("✅ Preprocessing complete.")
print(f"Feature matrix shape: {X.shape}")
print(f"Target distribution:\n{y.value_counts(normalize=True)}")
```

```
✅ Preprocessing complete.
Feature matrix shape: (6362620, 12)
Target distribution:
isFraud
0    0.998709
1    0.001291
Name: proportion, dtype: float64
```

#### 6. Stratified sampling

```
# 🎨 Stratified sampling: Keep all frauds, sample non-frauds
fraud_df = df_clean[df_clean['isFraud'] == 1]
non_fraud_df = df_clean[df_clean['isFraud'] == 0].sample(n=100000, random_state=42) # Adjust n if needed

# 🔄 Combine and shuffle
sampled_df = pd.concat([fraud_df, non_fraud_df]).sample(frac=1, random_state=42).reset_index(drop=True)

# 🎯 Split into features and target
X = sampled_df.drop('isFraud', axis=1)
y = sampled_df['isFraud']
print(f"✅ Sampled dataset shape: {X.shape}")
print(f"🎯 Sampled class distribution:\n{y.value_counts(normalize=True)}")
```

```
✅ Sampled dataset shape: (108213, 12)
🎯 Sampled class distribution:
isFraud
0    0.924103
1    0.075897
Name: proportion, dtype: float64
```

## 7. Train-test split

```
from sklearn.model_selection import train_test_split

# 🇺🇸 Stratified split for fair class balance in both sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

print("✅ Train-test split done.")
print(f"Train shape: {X_train.shape}, Test shape: {X_test.shape}")
```

```
↔️ ✅ Train-test split done.
Train shape: (86570, 12), Test shape: (21643, 12)
```

## 5. Feature scaling

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("✅ Feature scaling complete.")
```

```
↔️ ✅ Feature scaling complete.
```

## 6. Model Training + Evaluation (Random Forest Classifier)

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# 🇺🇸 Train Random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42, class_weight='balanced')
rf_model.fit(X_train_scaled, y_train)

# 🔍 Predict on test set
y_pred = rf_model.predict(X_test_scaled)
y_proba = rf_model.predict_proba(X_test_scaled)[:, 1] # For ROC-AUC

# 🇺🇸 Evaluation
print("📊 Classification Report:")
print(classification_report(y_test, y_pred))

print("📊 ROC AUC Score:", roc_auc_score(y_test, y_proba))


# 📊 Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Not Fraud", "Fraud"])
disp.plot(cmap="Blues")
plt.title("📊 Confusion Matrix")
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	20000
1	0.98	0.96	0.97	1643
accuracy			1.00	21643
macro avg	0.99	0.98	0.98	21643
weighted avg	1.00	1.00	1.00	21643

ROC AUC Score: 0.9990947352404138

/usr/local/lib/python3.11/dist-packages/IPython/core/pylabtools.py:151: UserWarning: Glyph 129534 (\N{RECEIPT}) missing from font(s)  
fig.canvas.print\_figure(bytes\_io, \*\*kw)



## 7. Applying XGBoost

```
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, ConfusionMatrixDisplay
```

```
# Calculate scale_pos_weight = (non-fraud cases / fraud cases) in train set
scale_pos_weight = (y_train == 0).sum() / (y_train == 1).sum()
print(f"📊 scale_pos_weight: {scale_pos_weight:.2f}")
```

```
# Initialize and train the model
xgb_model = XGBClassifier(
    n_estimators=100,
    max_depth=5,
    learning_rate=0.1,
    subsample=0.8,
    colsample_bytree=0.8,
    scale_pos_weight=scale_pos_weight,
    use_label_encoder=False,
    eval_metric='logloss',
    tree_method='hist' # Fastest for CPU or 'gpu_hist' if using GPU
)
```

```
xgb_model.fit(X_train_scaled, y_train)
print("✅ XGBoost model trained.")
```

📊 scale\_pos\_weight: 12.18  
/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [17:11:16] WARNING: /workspace/src/learner.cc:740: Parameters: { "use\_label\_encoder" } are not used.

```
warnings.warn(msg, UserWarning)
✅ XGBoost model trained.
```

## 8. Evaluate the Tuned Model

```
# Predict and evaluate
y_pred_xgb = xgb_model.predict(X_test_scaled)
y_proba_xgb = xgb_model.predict_proba(X_test_scaled)[:, 1]

print("📄 Classification Report (XGBoost):")
print(classification_report(y_test, y_pred_xgb))
```

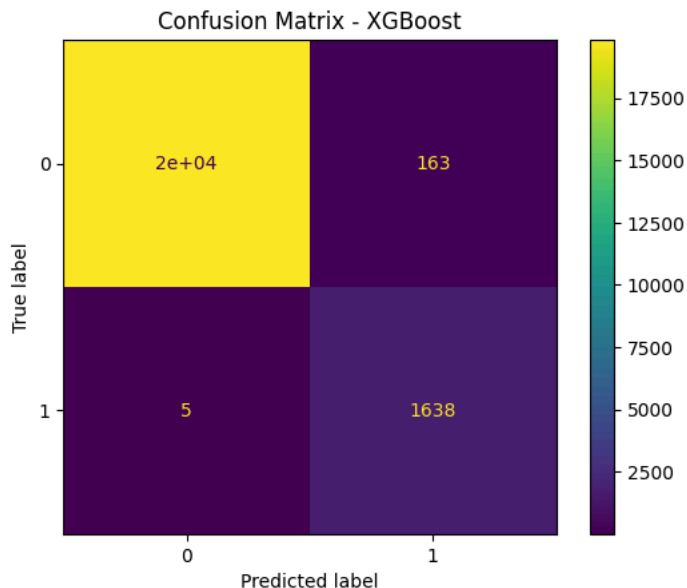
```
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_xgb)
```

```
plt.title("Confusion Matrix - XGBoost")
plt.show()

roc_auc_xgb = roc_auc_score(y_test, y_proba_xgb)
print(f"🔴 ROC-AUC Score (XGBoost): {roc_auc_xgb:.4f}")
```

```
🔄 📄 Classification Report (XGBoost):
```

	precision	recall	f1-score	support
0	1.00	0.99	1.00	20000
1	0.91	1.00	0.95	1643
accuracy			0.99	21643
macro avg	0.95	0.99	0.97	21643
weighted avg	0.99	0.99	0.99	21643



🔴 ROC-AUC Score (XGBoost): 0.9995

#### 9. Save the model

```
import joblib
from sklearn.preprocessing import StandardScaler

# Assuming you have already trained your XGBoost model
# For example: xgb_model.fit(X_train, y_train)

# If you are using a StandardScaler for scaling the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train) # Scale the training data

# Fit your model on the scaled data (replace with your actual model training step)
xgb_model.fit(X_train_scaled, y_train)

# Save the trained model to a .pkl file
model_path = "/content/drive/MyDrive/fraud_det/xgb_model.pkl"
joblib.dump(xgb_model, model_path)
print(f"✅ Model saved to: {model_path}")

# Save the scaler to a .pkl file
scaler_path = "/content/drive/MyDrive/fraud_det/scaler.pkl"
joblib.dump(scaler, scaler_path)
print(f"✅ Scaler saved to: {scaler_path}")
```

```
🔄 /usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [17:27:19] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.
```

```
warnings.warn(msg, UserWarning)
✅ Model saved to: /content/drive/MyDrive/fraud_det/xgb_model.pkl
✅ Scaler saved to: /content/drive/MyDrive/fraud_det/scaler.pkl
```

```
# Load the saved XGBoost model and scaler
xgb_model = joblib.load('/content/drive/MyDrive/fraud_det/xgb_model.pkl')
scaler = joblib.load('/content/drive/MyDrive/fraud_det/scaler.pkl')

# Now you can use the model and scaler to make predictions, as shown earlier
```

## 10. Running the model on a sample transaction

```
# Load model and scaler
import joblib
xgb_model = joblib.load("/content/drive/MyDrive/fraud_det/xgb_model.pkl")
scaler = joblib.load("/content/drive/MyDrive/fraud_det/scaler.pkl")

# Simulated data with correct order and columns
simulated_data = pd.DataFrame({
    'step': [50, 100],
    'amount': [250000, 1000000], # Very large amounts
    'oldbalanceOrig': [250000, 1000000], # Original sender balance
    'newbalanceOrig': [0, 0], # Balance drops to zero
    'oldbalanceDest': [0, 10], # Receiver had almost no balance
    'newbalanceDest': [250000, 1000010], # Receiver suddenly gains a large amount
    'isFlaggedFraud': [1, 1], # System thinks it's suspicious
    'type_CASH_IN': [0, 0],
    'type_CASH_OUT': [0, 1], # One is CASH_OUT
    'type_DEBIT': [0, 0],
    'type_PAYMENT': [0, 0],
    'type_TRANSFER': [1, 0] # Other is TRANSFER
})

# Scale simulated data
simulated_data_scaled = scaler.transform(simulated_data)

# Predict probabilities
probs = xgb_model.predict_proba(simulated_data_scaled)

# Apply custom threshold
threshold = 0.15
predictions = (probs[:, 1] >= threshold).astype(int)

# Display results
print("📊 Probabilities of fraud:", probs[:, 1])
print("🔴 Final Predictions with threshold:", predictions)
```


```
🔄 📊 Probabilities of fraud: [0.96058345 0.99930894]
🔴 Final Predictions with threshold: [1 1]
```

## 11. Visualization

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
sns.set(style="whitegrid")
%matplotlib inline
```


### Transaction type distribution

```
plt.figure(figsize=(8, 5))
sns.countplot(data=df, x='type', order=df['type'].value_counts().index, palette='muted')
plt.title('Transaction Type Distribution')
plt.xlabel('Transaction Type')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

 <ipython-input-36-c147d27fdda8>:2: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le

```
sns.countplot(data=df, x='type', order=df['type'].value_counts().index, palette='muted')
```




Fraud v/s Non-fraud count

```
plt.figure(figsize=(6, 4))
sns.countplot(data=df, x='isFraud', palette=['green', 'red'])
plt.title('Fraud vs Non-Fraud Transactions')
plt.xlabel('Is Fraud? (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.xticks([0, 1], ['Non-Fraud', 'Fraud'])
plt.tight_layout()
plt.show()
```

 <ipython-input-37-4877512cbb36>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le

```
sns.countplot(data=df, x='isFraud', palette=['green', 'red'])
```



Boxplot of Amounts by Fraud Label

```
plt.figure(figsize=(8, 5))
sns.boxplot(data=df, x='isFraud', y='amount', palette='coolwarm')
plt.yscale('log') # Amounts can be very large
plt.title('Transaction Amounts by Fraud Label')
plt.xlabel('Is Fraud?')
plt.ylabel('Transaction Amount (log scale)')
plt.tight_layout()
```

```
plt.show()
```

```
<ipython-input-38-148209421c8d>:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `le


```
sns.boxplot(data=df, x='isFraud', y='amount', palette='coolwarm')
```



Old vs New Balance (Orig Account) – Only for Frauds

```
fraud_df = df[df['isFraud'] == 1]
```

```
plt.figure(figsize=(8, 5))
sns.scatterplot(data=fraud_df, x='oldbalanceOrig', y='newbalanceOrig', hue='type', palette='Set2')
plt.title('Orig Account: Old vs New Balance (Frauds Only)')
plt.xlabel('Old Balance')
plt.ylabel('New Balance')
plt.tight_layout()
plt.show()
```



Correlation Heatmap

```
# Drop non-numeric columns before correlation
numeric_df = df.select_dtypes(include='number')
```

```
# Plot the correlation heatmap
```



```
plt.figure(figsize=(12, 8))
corr = numeric_df.corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f", square=True)
plt.title('Correlation Heatmap (Numerical Features Only)')
plt.tight_layout()
plt.show()
```

